

# Comp541 Final Project

## Gürsu Gülcü

“Dynamic Routing Between Capsules”  
by Sabour, Frosst, Hinton

# Problem 1: A CNN may identify both as a “person”. Why?



person

0.88



person

0.90



- It merely checks for the existence of some features it has learnt, not for the relationships among the features.

## Problem 2: A CNN fails to identify the inverted image.



person 0.88



reddish orange color 0.78



light brown color 0.78



starlet 0.66



entertainer 0.66



female 0.60



woman 0.59



young lady (heroine) 0.59



coal black color 0.79



hairpiece (hair) 0.71



dress 0.71



maroon color 0.71



person 0.58



toupee (hairpiece) 0.58



woman 0.56



Earrings 0.55




female 0.50



- Feature detection process is not as robust to transformations as us humans.
- Maybe our algorithms should try to mimic the human brain?


# Problem 3: Pooling operation may be discarding useful information.





- Sub-sampling due to pooling obtains crude spatial **invariance** by ignoring details -> Neural activities constant.
- Better to aim for **equivariance**: Invariance under transformations only.
  - Changes in viewpoint should change neural activities
  - Constant weights should code the viewpoint-invariant knowledge

# A Possible Solution: Inverse Computer Graphics

- Representation of objects in the brain does not depend on view angle
- Relationships between 3D objects can be represented by **pose**
  - **Translation** and **rotation**





*computer graphics: rendering engine*




*capsule network: inverse graphics*

A higher level entity is present if lower level visual entities can agree on their predictions for its pose




Predictions for face location from nose, mouth and eyes match -> there must be a face there

# Capsule Networks: A Solution to Our Problems?



- **Input** is the familiar 28x28x1 MNIST images.
- **First layer** is convolutional with 256 9x9 kernels and stride=1
  - No pooling!
  - Output is 20x20x256 since  $(28-9+1)/1 = 20$


# Second Convolutional Layer -> Primary Capsules



- **Second layer** is again convolutional with 256 9x9 kernels , stride=2
  - Output is 6x6x256 since  $(20-9+1)/2 = 6$




# Dissecting the Layer of Primary Capsules



- **Group 256 channels into 32**
  - Depth is 8
  - $6 \times 6 = 36$  localized “capsules”, 32 groups, each having a depth of 8
- **Interpretation:**
  - Z-axis (channels): Capsule contents
    - 32 features holding 8 numbers
    - Thickness, orientation etc.
  - X,Y axis (6x6 grid): Spatial info
    - Looking at a specific part of the image

# What are the Primary Capsules looking at?




- “Squash” the output vectors before passing to the next layer

# The Squashing Function

$$\mathbf{v}_j = \underbrace{\frac{\|\mathbf{s}_j\|^2}{1 + \|\mathbf{s}_j\|^2}}_{\text{additional "squashing"}} \underbrace{\frac{\mathbf{s}_j}{\|\mathbf{s}_j\|}}_{\text{unit scaling}}$$

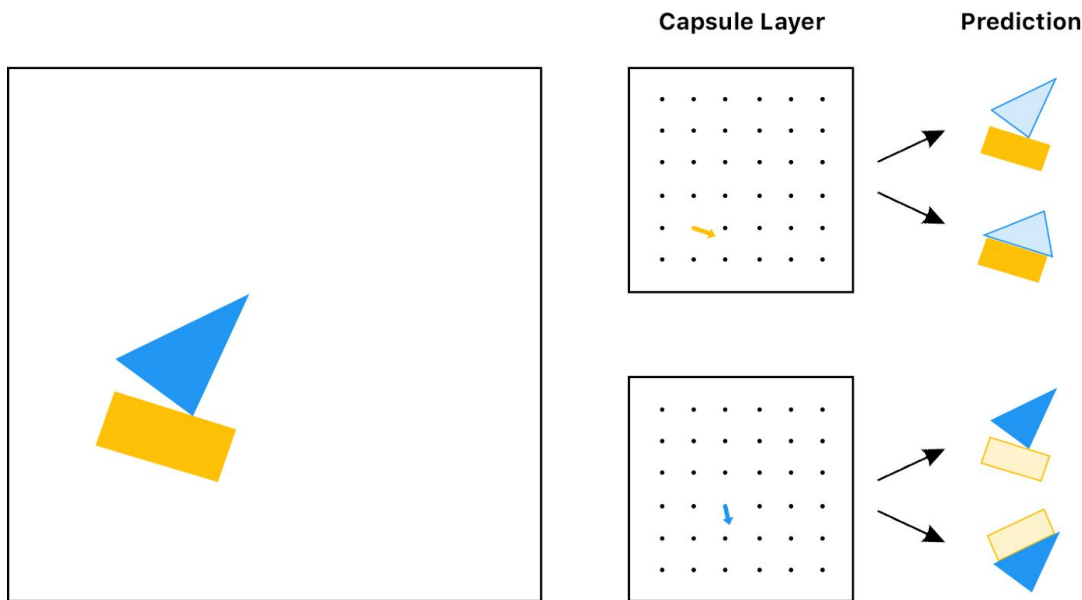
- Normalizing operation:
  - Unit vector for a long vector
  - Scale with its own small magnitude for further shrinkage for a short one

# Next Capsule Layer: Digit Caps



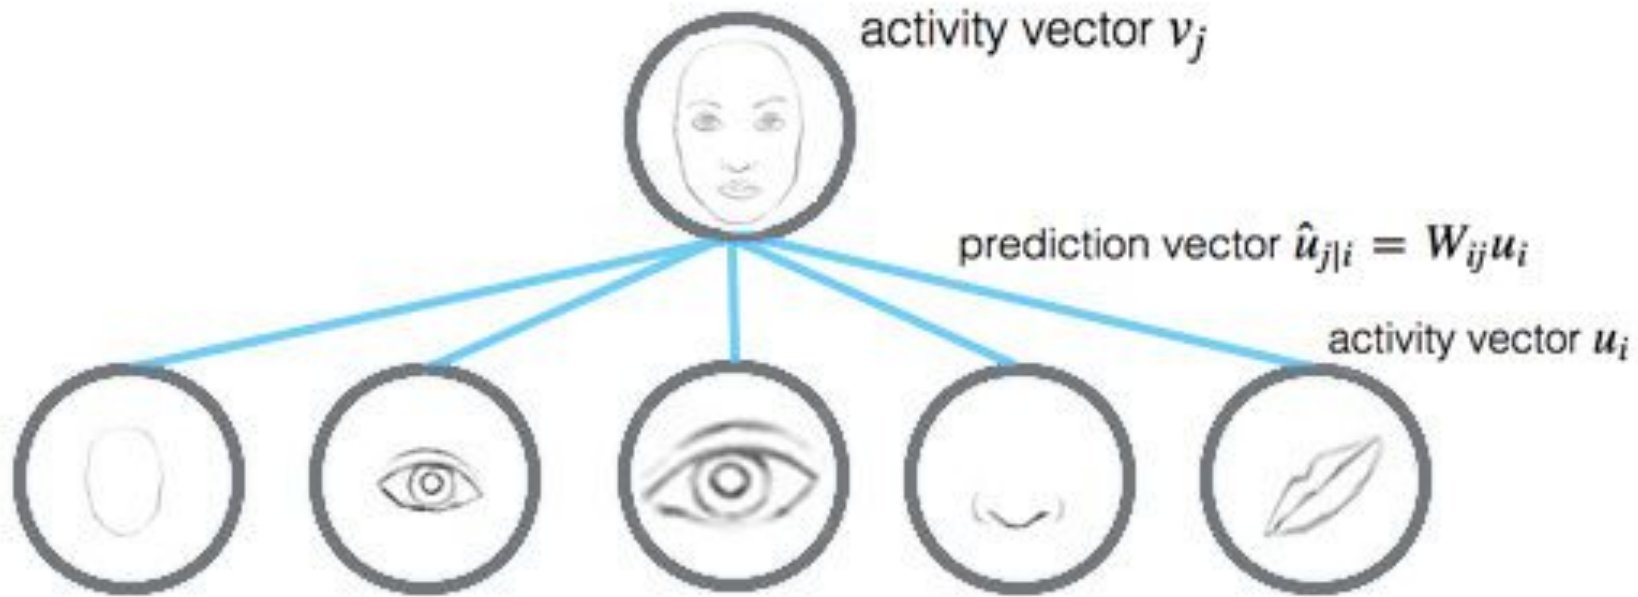
- **Digit Caps are 16-D**
  - Hold more information than Primary Caps (8-D) since they are looking at whole images
- An algorithm called **routing** is utilized to learn the weights of the transformation matrix  $W_{ij}$ 
  - Each  $W_{ij}$  is 8(in)x16(out)

# Role of Primary Capsule: “Which higher capsule am I a part of?”



- The square and the triangle will agree on the boat but disagree on the house due to the routing algorithm. How?

# Calculation for Predicting the Output of the Next Layer



$W_{ij}$  is  $16 \times 8$ ,  $u_i$  is  $8 \times 1$   $\rightarrow W_{ij}u_i = u_{j|i}$  is  $1 \times 16$   $\rightarrow$  Higher level capsules are 16-D

All are put in a big 1152 (Primary Capsules)  $\times$  10 (Digit Capsules) matrix

# Calculating the Affine Transformation Matrices is a challenge

$$\begin{pmatrix} \text{Batch size} \\ \begin{pmatrix} \mathbf{W}_{1,1} & \mathbf{W}_{1,2} & \cdots & \mathbf{W}_{1,10} \\ \mathbf{W}_{2,1} & \mathbf{W}_{2,2} & \cdots & \mathbf{W}_{2,10} \\ \vdots & \vdots & \ddots & \vdots \\ \mathbf{W}_{1152,1} & \mathbf{W}_{1152,2} & \cdots & \mathbf{W}_{1152,10} \end{pmatrix} \cdot \begin{pmatrix} \mathbf{u}_1 & \mathbf{u}_1 & \cdots & \mathbf{u}_1 \\ \mathbf{u}_2 & \mathbf{u}_2 & \cdots & \mathbf{u}_2 \\ \vdots & \vdots & \ddots & \vdots \\ \mathbf{u}_{1152} & \mathbf{u}_{1152} & \cdots & \mathbf{u}_{1152} \end{pmatrix} \end{pmatrix} = \begin{pmatrix} \begin{pmatrix} \hat{\mathbf{u}}_{111} & \hat{\mathbf{u}}_{211} & \cdots & \hat{\mathbf{u}}_{1011} \\ \hat{\mathbf{u}}_{112} & \hat{\mathbf{u}}_{212} & \cdots & \hat{\mathbf{u}}_{1012} \\ \vdots & \vdots & \ddots & \vdots \\ \hat{\mathbf{u}}_{111152} & \hat{\mathbf{u}}_{211152} & \cdots & \hat{\mathbf{u}}_{1011152} \end{pmatrix} \end{pmatrix}$$

- Multiplication of 1152x10xBatch\_Size pairs of matrices in each iter.
  - Note the dot, not a regular matrix multiplication of two huge matrices.
  - Must be done parallel on GPU, tile() & matmul() would come in handy.

# Details of the Routing Algorithm

---

**Procedure 1** Routing algorithm.

---

```
1: procedure ROUTING( $\hat{\mathbf{u}}_{j|i}, r, l$ )
2:   for all capsule  $i$  in layer  $l$  and capsule  $j$  in layer  $(l + 1)$ :  $b_{ij} \leftarrow 0$ .
3:   for  $r$  iterations do
4:     for all capsule  $i$  in layer  $l$ :  $\mathbf{c}_i \leftarrow \text{softmax}(\mathbf{b}_i)$  ▷ softmax computes Eq. 3
5:     for all capsule  $j$  in layer  $(l + 1)$ :  $\mathbf{s}_j \leftarrow \sum_i c_{ij} \hat{\mathbf{u}}_{j|i}$ 
6:     for all capsule  $j$  in layer  $(l + 1)$ :  $\mathbf{v}_j \leftarrow \text{squash}(\mathbf{s}_j)$  ▷ squash computes Eq. 1
7:     for all capsule  $i$  in layer  $l$  and capsule  $j$  in layer  $(l + 1)$ :  $b_{ij} \leftarrow b_{ij} + \hat{\mathbf{u}}_{j|i} \cdot \mathbf{v}_j$ 
   return  $\mathbf{v}_j$ 
```

---


- The  $r$  iterations increase the weight of primary  $u_i$  to estimate digit  $v_j$  when the affine transformation of  $u_i$  is aligned with  $v_j$  ( $\hat{\mathbf{u}}_{j|i} \cdot \mathbf{v}_j$ )
  - Repeatedly calculate the cluster center and adjust weights by cosine similarity:  $\cos(0)=1$  when aligned,  $\cos(90)=0$  when orthogonal
  - Would you agree that something is weird about **line 7**?




# Routing Algorithm In Action

(Darker=Bigger Value)


1




2



3



4



# Calculating the Loss

## CapsNet Loss Function

loss term for one DigitCap

calculated for correct DigitCap

calculated for incorrect DigitCaps

$$L_c = T_c \max(0, m^+ - ||\mathbf{v}_c||)^2 + \lambda (1 - T_c) \max(0, ||\mathbf{v}_c|| - m^-)^2$$

1 when correct DigitCap, 0 when incorrect

zero loss when correct prediction with probability greater than 0.9, non-zero otherwise

L2 norm

0.5 constant used for numerical stability

1 when incorrect DigitCap, 0 when correct


zero loss when incorrect prediction with probability less than 0.1, non-zero otherwise

L2 norm

Note: correct DigitCap is one that matches training label, for each training example there will be 1 correct and 9 incorrect DigitCaps


- Allows for detecting multiple images!
  - Can you detect a shortcoming?

# Decoder Architecture



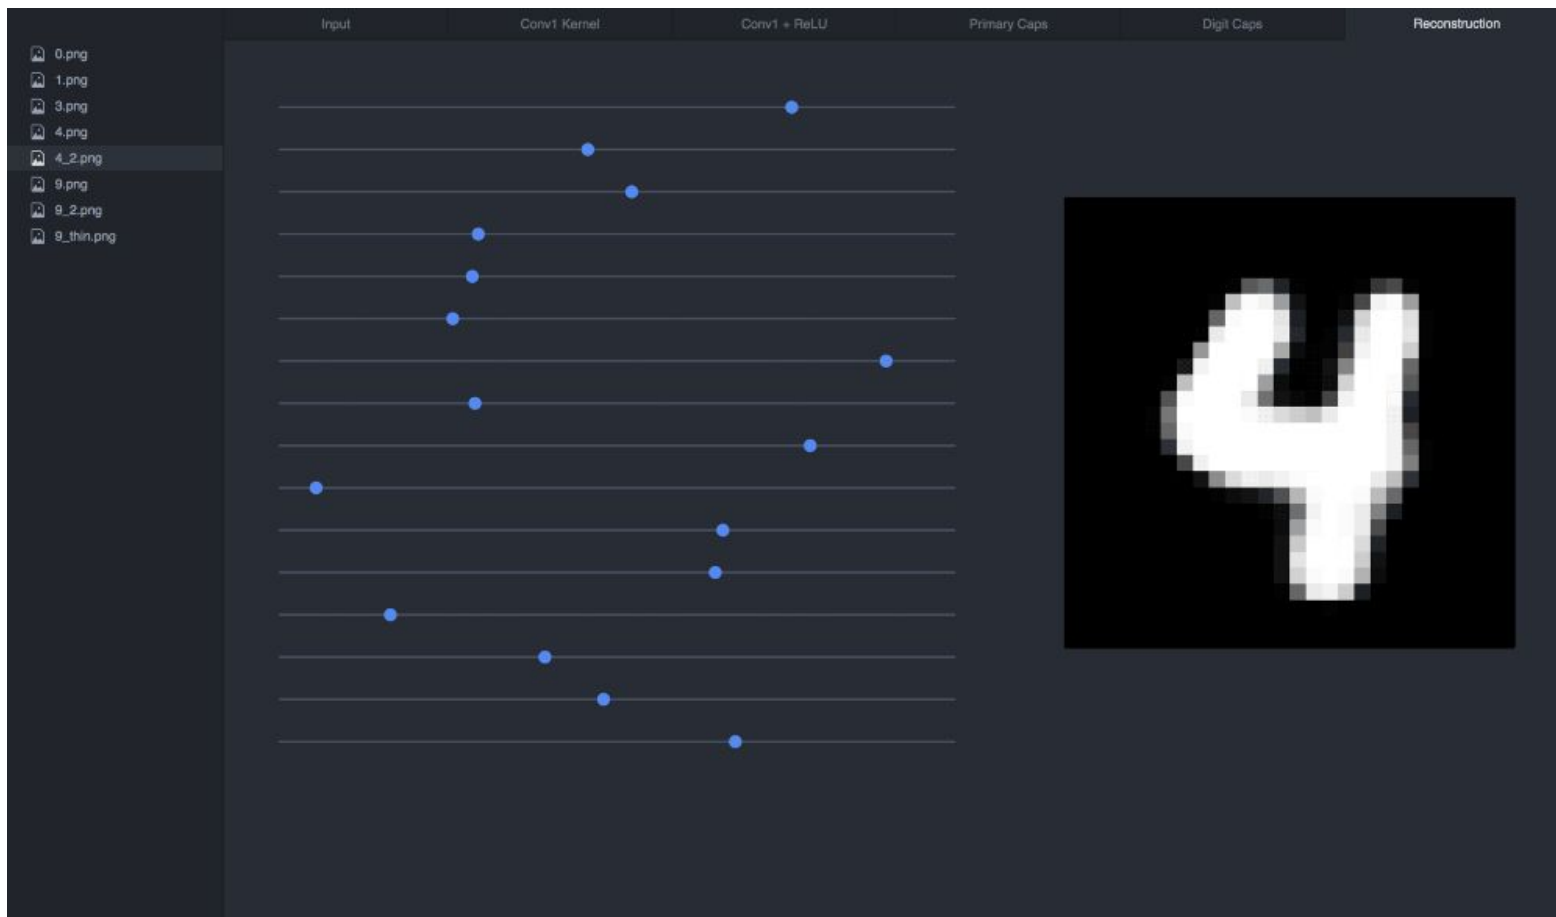
- Generates images from features to be compared against the input
  - The difference is an error, to be used as a regularizer

# Summary



- Squashing the output of each DigitCaps gives the prob. of existence
  - Do they add up to 1?

# Reconstruction Visualization





# MNIST Error Results by Methods: Table1




<u>METHOD</u>	<u>ROUTING r:</u>	<u>RECONSTRUCT</u>	<u>MNIST (%)</u>	<u>MY RESULT (Epoch)</u>
Baseline	-	-	0.39	0.56 (40)
CapsNet	1	NO	0.34	0.74 (5)
CapsNet	1	YES	0.29	0.61 (16)
CapsNet	3	NO	0.35	0.86 (9)
CapsNet	3	YES	0.25	0.65 (12)

# Promising Results for Viewpoint Invariance: The AffNIST Dataset



- In only 10 epochs ( $r=3$ ):
  - Train Error  $< 0.03\%$
  - Test Error  $< 2.5\%$
- No previous training on transformed images
- Used 60K out of 2M affine-transformed MNIST images

# Multi-MNIST Dataset: Not Publicly Available



- Tried to generate by superimposing AffNIST samples
- Remember sum of the output of Digit Caps  $\neq 1$ ?
  - Probability of existence for each digit -> Can detect multiple digits!
- Remember the shortcoming of Loss Function?
  - Can not count the number of occurrences of the same class
    - E.g. two 5's or one 5 in the picture?
- My initial results (>50%) not comparable to reported figures (5-8%)
  - Dataset, training proc., model parameters may be quite different



# Q & A